

Auditing Closed-Loop Learning in Recurrent Neural Networks: Reproduction, Robustness, and Generalization

Anonymous authors

Paper under double-blind review

Abstract

Recurrent neural networks are often used as mechanistic models of learning and control, but closed-loop training creates reproducibility challenges because the model’s actions alter future inputs. We conduct a claim-level reproducibility study of Ger and Barak’s analysis of closed-loop RNN learning dynamics. Rather than asking only whether original figures can be regenerated, we test whether the paper’s main claims survive independent implementation, seed variation, optimizer and hyperparameter perturbations, coupled-system stability diagnostics, and transfer to nonlinear or path-integration-style settings. Across 295 completed paired runs, the central open-loop/closed-loop divergence reproduces strongly: under the original double-integrator setting, 50/50 paired seeds show lower deployed loss for closed-loop training, with mean final deployed losses 0.0933 for closed-loop training versus 0.3119 for open-loop training. The coupled-system stability transition also reproduces in 50/50 original seeds. However, the stage claim is more nuanced: all original seeds show an early rapid-improvement phase followed by slow progress, but an objective three-stage changepoint criterion finds no robust second boundary. Generalization is partial: 38/50 architecture/task extension runs preserve deployed-loss divergence, while the simple ring path-integration variant does not. These results motivate concrete reporting practices for closed-loop RNN studies: deployed closed-loop loss, paired-seed comparisons, algorithmic stage criteria, coupled-system spectra, feedback strength, rollout horizon, and seed-level failure rates should be reported whenever papers make claims about recurrent learning dynamics in interactive environments.

1 Introduction

Closed-loop recurrent neural networks occupy an important middle ground between supervised sequence modeling, control, and computational neuroscience. In a closed-loop environment, an RNN’s output changes the future inputs it receives. This feedback makes the training problem qualitatively different from open-loop imitation or teacher-forced supervised learning, where future inputs are independent of the learner’s deployed actions. Ger and Barak (Ger and Barak, 2025) recently argued that this distinction produces learning dynamics that are fundamentally different from those of otherwise identical open-loop RNNs. Their paper is valuable because it makes a mechanistic claim rather than only a performance claim: the coupled agent-environment system, not the RNN alone, determines important learning transitions.

This paper is not a figure-level reproduction study. Instead, we use Ger and Barak’s closed-loop RNN learning framework as a test case for claim-level reproducibility in recurrent control systems. We ask which conclusions survive independent implementation, seed variation, optimizer and hyperparameter perturbations, coupled-system stability analysis, and transfer to nonlinear or path-integration-style settings. This framing matters for TMLR-style reproducibility: a rerun of a public artifact can be educational, but it is only broadly useful when it surfaces reusable lessons about when claims are reliable, when they are protocol-sensitive, and what future papers should report.

Our study makes four contributions. First, we provide an independent, config-driven reproduction of the central open-loop/closed-loop learning comparisons in Ger and Barak. Second, we convert qualitative claims about learning stages and stability transitions into quantitative, seed-level tests with confidence intervals. Third, we measure robustness under perturbations to optimizer, initialization, feedback strength, noise, and rollout horizon, identifying which conclusions are stable and which are protocol-sensitive. Fourth, we distill an audit checklist for future closed-loop RNN studies, including paired-seed comparisons, deployed closed-loop loss, stage criteria, coupled-system stability diagnostics, and failure-rate reporting.

2 Related Work

Task-trained RNNs in computational neuroscience. Task-trained RNNs are widely used as mechanistic models because they can solve cognitive and motor tasks while exposing internal dynamics for analysis. Early and influential work showed that high-dimensional trained RNNs can implement low-dimensional dynamical mechanisms (Sussillo and Barak, 2013), and that recurrent dynamics can account for context-dependent computation in prefrontal cortex (Mante et al., 2013). This line of work has become a general modeling strategy in systems neuroscience, where RNNs serve as hypothesis-generating models rather than only predictive tools (Barak, 2017). At the same time, large-scale analyses of trained recurrent networks show that architecture and training details can change representational geometry and dynamics even when task performance is similar (Maheswaranathan et al., 2019). This motivates our claim-level framing: reproducing a task curve is not enough if the mechanistic conclusion depends on training protocol or model class.

Dynamical-systems analysis of RNNs. The original paper sits in a tradition of reverse-engineering trained RNNs through fixed points, low-dimensional projections, eigenvalues, and linearized dynamics. Sussillo and Barak’s low-dimensional analysis made the case that trained recurrent networks can be understood through dynamical-systems tools (Sussillo and Barak, 2013), and later work used related analyses to compare recurrent dynamics across large populations of networks (Maheswaranathan et al., 2019). Our audit differs in emphasis. We do not only ask whether a trained RNN has interpretable internal dynamics; we ask whether the relevant dynamical object is the RNN alone or the coupled agent-environment system. This distinction is central in closed-loop settings because the environment transforms actions into future observations.

Open-loop training, closed-loop deployment, and distribution shift. The mismatch between teacher-forced training and deployed rollouts is well known in sequence prediction and imitation learning. Scheduled sampling was proposed to reduce the discrepancy between training on ground-truth previous tokens and testing on model-generated tokens (Bengio et al., 2015). In imitation learning, DAGger and related reductions explicitly address the fact that a learner’s actions change the future states it visits, invalidating i.i.d. supervised-learning assumptions (Ross et al., 2011). Closed-loop RNN learning has an analogous but mechanistically richer mismatch: open-loop imitation can minimize teacher-forced action error while still failing under deployed closed-loop rollouts. This connection motivates our insistence on reporting deployed closed-loop loss, not only open-loop or imitation loss.

Reproducibility and robustness in machine learning. Our audit follows the broader shift from artifact reruns to structured reproducibility studies. The NeurIPS reproducibility program emphasized clear artifacts, documented environments, and reproducibility checklists (Pineau et al., 2021). In deep reinforcement learning, Henderson et al. showed that algorithmic conclusions can be sensitive to seeds, evaluation protocol, and implementation details (Henderson et al., 2018). Closed-loop RNN studies share many of these risks because training outcomes can depend on random initialization, rollout horizon, optimizer, and environment coupling. The present paper therefore treats Ger and Barak’s work as a case study for a reusable audit protocol: paired seeds, stress tests, coupled-system diagnostics, and reporting lessons.

3 Original Claims and Audit Questions

We organize the study around five claims. C1 is the central divergence claim: identical RNNs trained open-loop and closed-loop follow different learning trajectories and differ under closed-loop deployment. C2 is the stage claim: closed-loop learning passes through distinct learning stages. C3 is the mechanistic stability

Table 1: Closed-loop RNN reproducibility audit protocol. The protocol is intentionally lightweight: each row defines a required audit step, its minimum implementation, and the output needed to support claim-level conclusions.

Audit step		Minimum requirement	Recorded output
Paired training		Clone the same initialization into open-loop and closed-loop learners; keep architecture and evaluation budget fixed.	Per-seed train loss, deployed closed-loop test loss, peak deployed loss, final loss, and effective-gain trajectory for both learners.
Seed-level	uncertainty	Run enough seeds to estimate direction agreement and confidence intervals; report seed count explicitly.	Mean, 95 percent bootstrap CI, fraction of seeds supporting each claim, and seed-level failure/recovery rates.
Stage detection		Predefine a detector using log-loss derivative, minimum plateau duration, and optional stability crossing.	Stage labels, plateau duration, plateau loss, transition epochs, and frequency of detected stages across seeds/settings.
Coupled diagnostics		Compute RNN-only and coupled agent-environment spectra whenever the system can be linearized or locally linearized.	Spectral radius time series, stability-crossing epoch, plateau-exit gap, and predictive correlation with recovery/failure.
Robustness perturbations		Change factors that alter optimization or environment coupling, such as optimizer, initialization scale, feedback strength, noise, horizon, control penalty, and clipping.	Claim support under each perturbation, effect-size changes, and perturbation-level failure rates.
Generalization check		Test at least one architecture or task variant that preserves action-dependent inputs and one that changes the recurrent/control structure.	Claim support by variant and a boundary-condition statement explaining when the phenomenon does or does not transfer.
Claim decision		Classify each claim as reproduced, partially reproduced, or not reproduced using the same criteria across settings.	Claim-level matrix linking evidence type, original setting, robustness, generalization, and actionable lesson.

claim: progress depends on stability of the coupled agent-environment system. C4 is the protocol robustness claim: short-term policy improvement and long-term stability interact with optimizer, horizon, feedback strength, and other choices. C5 is the generalization claim: similar dynamics appear beyond the simplest double-integrator setting.

For each claim, we define three evidence levels. A claim is *reproduced* when the direction of the effect and the relevant diagnostic hold under the original protocol for most paired seeds with confidence intervals excluding negligible effects. A claim is *partially reproduced* when the direction appears under the original protocol but is unstable across seeds, sensitive to plausible perturbations, or only weakly supported by the diagnostic. A claim is *not reproduced* when the effect is absent, reverses, or depends on hand-picked visual criteria. These labels are intended to avoid the ambiguous statement that a result “matches the original” without specifying the evidence.

4 Closed-Loop RNN Reproducibility Audit Protocol

We specify a modest audit protocol rather than proposing a new benchmark. The protocol is a recipe for testing claims about closed-loop RNN learning dynamics in any recurrent control setting where actions can affect future observations. Its inputs are: a set of original claims, a task/environment, an open-loop training procedure, a closed-loop training procedure, a seed set, a small set of protocol perturbations, and at least one held-out architecture or task variant. Its outputs are: per-seed deployed closed-loop metrics, stage labels, stability diagnostics, perturbation summaries, generalization summaries, and a claim-level reproducibility matrix. Table 1 gives the concrete specification used in this paper.

The protocol deliberately separates *measurement* from *interpretation*. Measurement consists of paired runs, deployed losses, spectra, stage labels, and perturbation outcomes. Interpretation happens only after aggre-

gating seed-level outputs into claim-level decisions. This separation is important because closed-loop RNN papers often make qualitative statements about stages or mechanisms from a small number of curves; our protocol requires those statements to be tied to explicit detectors, uncertainty estimates, and failure rates.

4.1 Exact metric definitions

Let $s_t \in \mathbb{R}^{d_s}$ be the environment state, $o_t = Cs_t + \epsilon_t$ the observation, h_t the RNN hidden state, and $u_t = f_\theta(o_t, h_t)$ the action. Closed-loop deployment follows

$$s_{t+1} = F(s_t, u_t), \quad h_{t+1} = G_\theta(o_t, h_t). \quad (1)$$

For an episode of horizon T , the deployed closed-loop loss for model θ is

$$L_{\text{deploy}}(\theta) = \frac{1}{T} \sum_{t=0}^{T-1} [\ell_s(s_t, s_t^*) + \lambda_u \|u_t\|_2^2], \quad (2)$$

where ℓ_s is squared state or tracking error and λ_u is the control penalty. For the double-integrator stabilization task, $\ell_s(s_t, 0) = \|s_t\|_2^2$.

Open-loop imitation is evaluated separately from deployed performance. Given a teacher policy θ^* and teacher-generated observations o_t^* , the teacher-forced imitation loss is

$$L_{\text{TF}}(\theta; \theta^*) = \frac{1}{T} \sum_{t=0}^{T-1} \|f_\theta(o_t^*, h_t) - f_{\theta^*}(o_t^*, h_t^*)\|_2^2. \quad (3)$$

The distinction between L_{TF} and L_{deploy} is central: C1 is supported only if the training-regime comparison is evaluated under deployed closed-loop rollouts.

To summarize the controller in the double-integrator setting, we estimate an effective linear feedback gain K_t from rollout states and actions. For a collection of state-action pairs (x_i, u_i) at training epoch t , with matrices $X = [x_1^\top; \dots; x_n^\top]$ and $U = [u_1^\top; \dots; u_n^\top]$, the ridge-stabilized estimate is

$$\hat{K}_t = (X^\top X + \alpha I)^{-1} X^\top U, \quad (4)$$

reported as a row vector for scalar control. The closed/open gain-trajectory distance is

$$D_K = \left(\sum_{t \in \mathcal{E}} \left\| \hat{K}_t^{\text{CL}} - \hat{K}_t^{\text{OL}} \right\|_F^2 \right)^{1/2}, \quad (5)$$

where $\mathcal{E}$ is the set of recorded epochs. In the current implementation every epoch is recorded; future larger sweeps can subsample $\mathcal{E}$ without changing the definition.

For linear RNN controllers and linear environments, we compute both RNN-only and coupled-system spectra. With

$$s_{t+1} = As_t + Bu_t, \quad h_{t+1} = (1 - \eta)h_t + \eta(W_{hh}h_t + W_{ih}Cs_{t+1}), \quad u_t = W_{ho}h_t, \quad (6)$$

the coupled transition matrix over $z_t = [s_t^\top, h_t^\top]^\top$ is

$$P_\theta = \begin{bmatrix} A & BW_{ho} \\ \eta W_{ih}CA & (1 - \eta)I + \eta W_{hh} + \eta W_{ih}CBW_{ho} \end{bmatrix}. \quad (7)$$

The two spectral diagnostics are

$$\rho_{\text{RNN}}(\theta) = \max_i |\lambda_i(W_{hh})|, \quad \rho_{\text{coup}}(\theta) = \max_i |\lambda_i(P_\theta)|. \quad (8)$$

We define the stability-crossing epoch as

$$\tau_\rho = \min \{t : \rho_{\text{coup}}(\theta_t) < 1 \text{ and } \rho_{\text{coup}}(\theta_{t:t+m}) < 1\}, \quad (9)$$

where m is a minimum persistence window. The plateau-exit gap is $\Delta_\rho = \tau_{\text{exit}} - \tau_\rho$.

Stage labels are computed from smoothed log deployed loss $y_t = \log(\max(L_{\text{deploy},t}, \varepsilon))$. Let $\dot{y}_t$ be a finite-difference derivative. Stage 1 ends at the first epoch τ_1 such that

$$|\dot{y}_{t:t+m}| \leq \delta_{\text{plat}} \quad (10)$$

for at least m consecutive epochs. The plateau exits at the first later epoch τ_2 such that

$$\dot{y}_t < -\delta_{\text{improve}}, \quad (11)$$

or at the stability crossing τ_ρ when no derivative-based exit is detected. The plateau length is $\tau_2 - \tau_1$, and C2 support is the fraction of seeds/settings with plateau length at least m .

For a scalar metric M and paired seed i , the open/closed effect is

$$\Delta_i(M) = M_i^{\text{OL}} - M_i^{\text{CL}}. \quad (12)$$

We report the mean effect $\bar{\Delta}$, a nonparametric bootstrap 95 percent confidence interval, and the direction-agreement fraction

$$a(M) = \frac{1}{N} \sum_{i=1}^N \mathbf{1}[\Delta_i(M) > 0]. \quad (13)$$

For a claim c with binary seed-level support indicator $q_i(c)$, the claim-support rate is

$$S(c) = \frac{1}{N} \sum_{i=1}^N q_i(c). \quad (14)$$

Failure rates are defined analogously by replacing $q_i(c)$ with indicators for divergence, non-finite loss, missing plateau, or failure to recover by the final epoch.

4.2 Paired open-loop and closed-loop training

For each seed, we initialize one recurrent controller and clone its weights. The closed-loop model is trained by rolling the controller in the environment and minimizing deployed state cost plus a control penalty. The open-loop model is trained as a student to imitate a teacher trajectory generated by the trained closed-loop controller. Both models are then evaluated under closed-loop deployment. The key dependent variables are final deployed loss, peak deployed loss, effective feedback-gain trajectory, failure or recovery status, and time to stable performance.

4.3 Stage detection

The original paper describes closed-loop learning stages visually. We operationalize this claim. Stage 1 is early rapid improvement, Stage 2 is a plateau or slow-progress phase, and Stage 3 is post-transition improvement. The detector uses log-loss derivatives, a minimum plateau duration, and optionally a coupled spectral-radius crossing. This makes stage claims falsifiable: a stage either appears consistently across seeds and perturbations or it does not.

4.4 Coupled-system stability diagnostics

For linearized settings, we compute the spectral radius of the RNN recurrent matrix and of the coupled agent-environment transition matrix. The audit asks whether the coupled-system spectral radius predicts plateau exit, instability, or recovery better than the RNN-only spectral radius. This test directly targets the mechanistic contribution of the original paper: if closed-loop dynamics are genuinely about the agent-environment loop, the coupled diagnostic should be more predictive than an isolated RNN diagnostic.

4.5 Robustness and generalization

The robustness sweep changes learning rate, initialization scale, feedback strength, episode length, observation noise, control penalty, optimizer, and gradient clipping. The generalization sweep tests nonlinear RNNs, GRUs, low-rank RNNs, a tracking-control variant, and a ring/path-integration-style task. We do not treat every extension as a new benchmark. Instead, each extension asks whether the closed-loop/open-loop distinction persists when the task still has action-dependent inputs and partial observability.

5 Result 1: Direct Reproduction of the Core Divergence

The central open-loop/closed-loop divergence reproduces cleanly under the original double-integrator setting. Across 50 paired seeds, the final deployed closed-loop loss of the closed-loop trained model is 0.0933 with bootstrap 95 percent CI [0.0932, 0.0935]. The open-loop trained student, evaluated under deployed closed-loop rollout, has final loss 0.3119 with CI [0.3044, 0.3193]. The paired deployed-loss gap is therefore 0.2186 with CI [0.2110, 0.2260], and the direction of the effect is reproduced in 50/50 seeds. Effective feedback gains also diverge: the final closed/open gain distance is 0.5045 with CI [0.4998, 0.5092].

The result is not just a visual overlay of curves. It shows that minimizing open-loop imitation loss can produce a controller whose deployed closed-loop loss remains substantially higher than the directly closed-loop trained controller. This supports C1 and the main reporting lesson of the audit: closed-loop RNN studies should report deployed closed-loop performance, not only teacher-forced or open-loop training loss.

6 Result 2: Which Claims are Robust?

The robustness sweep consists of 195 paired runs: 13 perturbation settings with 15 seeds per setting. C1 is robust across most perturbations. It is supported in 165/195 robustness runs overall and in every finite run for the original, learning-rate, initialization, feedback-strength, episode-length, observation-noise, and high-control-penalty perturbations. The two important exceptions are informative. With Adam, the final deployed loss gap almost disappears: mean final closed-loop loss is 0.0780 and mean final open-loop deployed loss is 0.0798, yielding no C1 support under the 5 percent deployed-loss criterion. With gradient clipping disabled, all 15 runs in the `no_clip` condition produce non-finite final losses and are counted as failures rather than negative evidence for the mechanistic claim.

The stage claim is more sensitive to how “stage” is defined. A slow-progress phase is detected in 180/195 robustness runs, with failures concentrated entirely in the non-finite no-clipping condition. However, the stricter three-stage criterion, which requires both a plateau and a non-fallback second boundary, is supported in only 1/195 robustness runs. Thus the broad qualitative observation that closed-loop training can enter a slow-progress regime is robust, but the claim that three distinct stage boundaries are algorithmically stable is not supported by the current detector.

The coupled-system stability transition is robust wherever the coupled spectrum is available and the run remains finite. It is supported in 180/195 robustness runs, again failing only in the no-clipping condition. This supports the original mechanistic emphasis on the coupled agent-environment system and separates the stability claim from the more fragile stage-boundary claim.

The current C4 evidence should be treated as incomplete. A weak proxy—open-loop deployed-loss spike together with closed-loop recovery—is supported in 32/295 completed runs and in 0/50 original-setting runs. This proxy is not the same as the paper’s short-term versus long-term stability tradeoff. We therefore do not use it to reject C4; instead, it motivates the targeted multi-horizon control-penalty sweep included in the artifact for the final audit.

7 Result 3: Which Diagnostics are Predictive?

The original-setting coupled-system stability transition reproduces in 50/50 seeds. The final coupled spectral radius is finite and below the stability threshold in every original run, and the stability-crossing epoch occurs

at the same point identified by the derivative-based plateau-start detector. This supports C3 in the original double-integrator setting.

However, the spectral diagnostic is architecture-dependent in the current implementation. For linear, tanh, and low-rank controllers on linear tasks, the coupled transition matrix can be constructed explicitly. For GRU and ring variants, the current audit reports `nan` for the coupled spectrum because no local linearization of the gated recurrent dynamics is implemented. This is not a failure of the original claim, but it is a practical limitation of the audit: coupled-system spectral claims require either an explicitly linearizable controller/environment pair or an additional local-linearization procedure.

The most important diagnostic distinction concerns C2. The coupled stability crossing is robust, but the strict three-stage plateau-exit boundary is not. In the original setting, a slow-progress phase is present in 50/50 seeds, yet the second boundary is a fallback in 50/50 seeds. A segmented-regression changepoint analysis likewise finds two log-loss segments that fit substantially better than one segment, but the third segment does not show the required post-plateau reacceleration. The median changepoint boundaries are epochs 40 and 255, with median slopes -0.1131 , -0.00229 , and -0.00036 for the three fitted segments. Thus our data support “rapid improvement followed by slow progress,” but not a robust three-stage sequence with a clearly detected exit.

8 Result 4: Does the Phenomenon Generalize?

Generalization is partial. Across 50 completed architecture/task extension runs, C1 deployed-loss divergence is supported in 38/50 runs. The effect transfers cleanly to the tanh RNN, GRU, and low-rank variants: each shows C1 support in 10/10 seeds. The tracking task shows partial support, with deployed-loss divergence in 8/10 seeds. The simple ring path-integration variant does not support C1 under the current task definition: 0/10 seeds show deployed-loss divergence and the mean deployed-loss gap is only 5.4×10^{-4} .

The negative ring result is interpretable as a boundary condition rather than a direct contradiction of the original double-integrator claim. The current ring task gives the learner $\sin(\text{error})$, $\cos(\text{error})$, and velocity, with a fixed target. This makes the deployed state-action mapping comparatively easy for the open-loop student to imitate. In contrast, the variants that preserve stronger closed-loop mismatch or harder control structure show the divergence. The artifact therefore includes a harder path-integration follow-up with partial observation and moving targets; those results are not included in the 50-run generalization numbers reported here.

C3 generalizes only where the coupled spectrum is implemented. It is supported in 10/10 tanh RNN runs and 10/10 low-rank runs, but not reported for GRU or ring runs because the current spectral code does not linearize gated nonlinear controllers. This is another reporting lesson: a generalization claim about stability mechanisms should distinguish performance transfer from diagnostic transfer.

9 Claim-Level Reproducibility Matrix

Table 2 is the main summary of the paper. It is designed to make the audit more than a class-report rerun: each original claim is attached to evidence, robustness, generalization, and an actionable lesson.

10 Actionable Lessons for Closed-Loop RNN Studies

The goal of the audit is to produce reusable reporting practices. We propose the following checklist for closed-loop RNN papers.

Report deployed closed-loop performance, not only open-loop or teacher-forced loss. The central risk is that a model can imitate a teacher under fixed inputs but fail when its own actions determine future inputs.

Table 2: Claim-level reproducibility matrix for completed runs. C4 is marked incomplete because the completed audit contains only a weak proxy; the artifact includes a targeted multi-horizon tradeoff sweep for the final version.

Claim	Evidence type	Original protocol	Robustness sweep	Generalization	Lesson
C1: open/closed divergence	Paired seeds, effect size, CI, deployed loss	Reproduced: 50/50 seeds, loss gap 0.2186	Robust in most perturbations; absent under Adam and failed without clipping	38/50 extension runs; fails on simple ring task	Report deployed closed-loop loss, not only teacher-forced loss.
C2: stage structure	Stage detector, plateau duration, changepoints	Partial: slow-progress phase 50/50, strict three-stage 0/50	Slow phase robust; strict exit boundary fragile	Slow phase common; strict boundary varies by task	Stage claims require algorithmic criteria, not visual identification.
C3: coupled stability	Coupled spectral radius crossing	Reproduced: 50/50 seeds	Robust in all finite linearizable robustness settings	Holds for tanh/low-rank; not implemented for GRU/ring	Report coupled-system spectrum, not only RNN-only spectrum.
C4: short/long tradeoff	Multi-horizon tradeoff metric	Incomplete; proxy 0/50	Proxy mixed and not decisive	N/A	Use targeted horizon/control-penalty sweeps for tradeoff claims.
C5: broader generalization	Architecture and task extensions	N/A	N/A	Partial: tanh/GRU/low-rank yes, tracking 8/10, ring 0/10	Closed-loop effects depend on environment coupling and partial observability.

Use paired-seed comparisons when contrasting open-loop and closed-loop training. The clean comparison uses the same initialization, architecture, optimizer budget, and evaluation rollouts while changing only the training regime.

Define learning stages algorithmically. If a paper claims stages, it should specify transition criteria, minimum plateau duration, and how often each stage appears across seeds.

Analyze the coupled agent-environment system. RNN-only spectra can miss stability changes introduced by the environment. Coupled diagnostics should be reported whenever a mechanistic stability claim is made.

Report feedback strength and rollout horizon. These parameters determine how strongly actions affect future observations and can control whether the closed-loop phenomenon appears.

Report failure rates, not just average curves. Means can hide unstable runs, recovery failures, or seed-dependent plateaus.

Include at least one stress test. Optimizer, initialization, noise, feedback strength, and horizon perturbations are inexpensive compared with the interpretive cost of an untested mechanistic claim.

11 Discussion

This project treats Ger and Barak’s framework as a case study in reproducibility for interactive recurrent systems. The completed audit supports the two strongest claims in the original setting: closed-loop and open-loop training diverge under deployment, and the coupled agent-environment system undergoes a stability transition. It also identifies two important limits. First, the stage claim depends strongly on how stages are defined. A slow-progress regime is reproducible, but a robust three-stage sequence with an objective second

boundary is not recovered by our detectors. Second, generalization is conditional: the divergence transfers across several architectures and a tracking-control task, but disappears in a simple ring task that gives the open-loop student enough state information to imitate the deployed policy well.

These mixed results are the reason a claim-level audit is useful. A figure-level reproduction would likely emphasize that the main closed-loop/open-loop loss gap appears. A methodological reproducibility study can say more: optimizer choice can remove the deployed-loss gap, missing gradient clipping can produce outright failure, spectra are only available for linearizable controllers unless local linearization is added, and stage claims should be accompanied by algorithmic boundaries rather than visual annotations alone. This is especially important for computational neuroscience, where RNNs are often used as mechanistic models and where closed-loop behavior is closer to biological learning than fixed-dataset supervision.

12 Limitations

The audit does not cover every RNN architecture, every control task, or full reinforcement-learning algorithms. The current path-integration extension is intentionally simple and, as the negative ring result shows, may be too easy to stress the open-loop/deployed mismatch. A harder partial-observation moving-target ring task is included in the artifact as a targeted follow-up but is not included in the completed-run numbers reported above. The C4 short-term versus long-term tradeoff is also not settled by the completed runs because our completed metric is only a weak proxy; the repository now includes a multi-horizon control-penalty sweep for a direct test. Finally, coupled spectral analysis is implemented for linearizable controllers and environments, but not yet for GRUs or other gated nonlinear controllers.

13 Artifact

The accompanying repository provides a config-driven implementation with smoke and full-run modes, raw CSV outputs, claim tables, and figure-generation scripts. The artifact is organized around the audit protocol rather than notebook-only reproduction: each experiment records the config, seed, device, metrics, and per-epoch time series. Full-run results can be regenerated with `bash scripts/run_full_audit.sh`; a lightweight validation run is available with `bash scripts/run_smoke.sh`. The targeted follow-up experiments for C2, C4, and the harder C5 path-integration setting can be launched with `bash scripts/run_targeted_c2_c4_c5.sh`.

References

- Yoav Ger and Omri Barak. Learning dynamics of RNNs in closed-loop environments. In *Advances in Neural Information Processing Systems*, poster, 2025. OpenReview: <https://openreview.net/forum?id=P63bBMXCiH>.
- Omri Barak. Recurrent neural networks as versatile tools of neuroscience research. *Current Opinion in Neurobiology*, 46:1–6, 2017.
- Samy Bengio, Oriol Vinyals, Navdeep Jaitly, and Noam Shazeer. Scheduled sampling for sequence prediction with recurrent neural networks. In *Advances in Neural Information Processing Systems*, 2015.
- Peter Henderson, Riashat Islam, Philip Bachman, Joelle Pineau, Doina Precup, and David Meger. Deep reinforcement learning that matters. In *Proceedings of the AAAI Conference on Artificial Intelligence*, 2018.
- Niru Maheswaranathan, Alex H. Williams, Matthew D. Golub, Surya Ganguli, and David Sussillo. Universality and individuality in neural dynamics across large populations of recurrent networks. In *Advances in Neural Information Processing Systems*, 2019.
- Valerio Mante, David Sussillo, Krishna V. Shenoy, and William T. Newsome. Context-dependent computation by recurrent dynamics in prefrontal cortex. *Nature*, 503:78–84, 2013.

Joelle Pineau, Philippe Vincent-Lamarre, Koustuv Sinha, Vincent Larivière, Alina Beygelzimer, Florence d’Alché-Buc, Emily Fox, and Hugo Larochelle. Improving reproducibility in machine learning research: a report from the NeurIPS 2019 reproducibility program. *Journal of Machine Learning Research*, 22(164):1–20, 2021.

Stéphane Ross, Geoffrey Gordon, and Drew Bagnell. A reduction of imitation learning and structured prediction to no-regret online learning. In *Proceedings of the Fourteenth International Conference on Artificial Intelligence and Statistics*, pages 627–635, 2011.

David Sussillo and Omri Barak. Opening the black box: low-dimensional dynamics in high-dimensional recurrent neural networks. *Neural Computation*, 25(3):626–649, 2013.

Ian Goodfellow, Yoshua Bengio, and Aaron Courville. *Deep Learning*. MIT Press, 2016.